

Naveed Mohiuddin

Data Engineer | AWS & Azure | Databricks

✉ navedmohiuddin0311@gmail.com | ☎ +1 773-394-9716

🌐 [linkedin.com/in/naveed-mohiuddin](https://www.linkedin.com/in/naveed-mohiuddin) | 🏠 github.com/NaveedMohiuddin | 🌐 navedmohiuddin.vercel.app | 📍 United States

PROFESSIONAL SUMMARY

Data Engineer with hands-on experience building batch and incremental pipelines across AWS, Azure, and Databricks. Work includes S3, Lambda, Glue, Athena, Airflow, Azure Data Factory, ADLS Gen2, Databricks, Delta Lake, PySpark, and Spark SQL, with an emphasis on reliable reruns, data quality, production troubleshooting, and analytics-ready datasets. Holds a Master of Science in Computer Science and AWS Solutions Architect, AWS Data Engineer, and Azure Fundamentals certifications.

TECHNICAL SKILLS

Cloud Platform: AWS (S3, Glue, Glue Data Catalog, Lambda, Redshift, Athena, EventBridge, Kinesis, SNS, EC2, RDS, Lake Formation), Azure (Data Factory, Databricks, Synapse Analytics, ADLS Gen2, Unity Catalog, Azure DevOps)
Data Engineering: ETL/ELT, Data-Lake & Medallion Architecture, Dimensional Modeling, SCD Type 2, Batch / Incremental / Streaming Pipelines, Data Quality, Data Governance & Lineage, Schema Evolution, Reconciliation
Big Data & Lakehouse: Apache Spark, PySpark, Spark SQL, Spark Structured Streaming, Delta Lake, Delta Live Tables (Lakeflow), Auto Loader, Databricks Workflows, Kafka, Hadoop, HDFS, Hive, MapReduce, Parquet
Databases & Warehouses: Amazon Redshift, Athena, Azure Synapse Analytics, Snowflake, SQL Server, Oracle, PostgreSQL, MySQL
Orchestration: Apache Airflow / MWAA (DAGs, SLA monitoring, retries), Amazon EventBridge, Azure Data Factory, Databricks Workflows
DevOps & Version Control: Git, GitHub, Bitbucket, Azure DevOps, GitHub Actions, Jenkins, CI/CD, Docker, Linux
Programming: Python, SQL, PySpark, Shell Scripting (Bash), Java, Pandas
Analytics & BI: Amazon QuickSight, Power BI, Athena SQL, Databricks SQL
AI-Assisted Development: GitHub Copilot, Claude (pipeline prototyping, data-quality automation)

CERTIFICATIONS

AWS Certified Solutions Architect – Associate
AWS Certified Data Engineer – Associate
Microsoft Azure Fundamentals

EXPERIENCE

Data Engineer

Benda Infotech

Jan 2025 – Present

Remote, US

- Built cloud data pipelines across AWS, Azure, and Databricks for batch ingestion, transformation, validation, and reporting use cases.
- On AWS, developed Lambda and EventBridge-based ingestion into S3, with raw and curated zones designed for traceability and targeted reprocessing.
- On Azure, built ADF and Databricks pipelines using Delta Lake medallion layers for incremental ingestion, transformation, and curated reporting.
- Developed PySpark and Spark SQL logic for data cleaning, validation, joins, reconciliations, and curated table creation.
- Used Delta **MERGE INTO** and deterministic S3 object naming to make reruns safe and reduce duplicate-load issues.
- Tuned Athena and Delta workloads using Parquet, partitioning, partition pruning, **OPTIMIZE**, and **Z-ORDER**.
- Added record-count, load-completion, and quarantine checks so failed or incomplete data could be caught before reaching reporting layers.
- Managed orchestration through Airflow, MWAA, ADF, and Databricks Workflows with retries, alerts, dependencies, and backfill support.
- Built Athena, Databricks SQL, QuickSight, and Power BI-ready datasets for operational and business reporting.
- Worked with analysts and engineers to translate reporting requirements into source mappings, validation rules, and reusable data models.
- **Tech:** AWS S3, Lambda, EventBridge, Glue, Athena, MWAA, Airflow, Azure Data Factory, Azure Databricks, Delta Lake, ADLS Gen2, PySpark, Spark SQL, Databricks SQL, Parquet, QuickSight, Power BI, Python, SQL, Git, Azure DevOps

Data Engineer

Applied Information Sciences

Feb 2021 – Jul 2023

Hyderabad, India

- Built and supported Azure-based data pipelines for policy, claims, and operational data across ADF, Databricks, Snowflake, and SQL-based sources.
- Developed ADF ingestion pipelines from SQL Server, Oracle, REST APIs, and flat files into ADLS Gen2 and Snowflake.
- Used control-table driven pipeline patterns to reduce repeated ADF development and make new-source onboarding more configuration-based.
- Implemented full-load and incremental extraction using watermark columns, runtime parameters, and validation checkpoints.
- Wrote PySpark, Spark SQL, and SQL transformations to standardize schemas, apply validation rules, and prepare curated data.
- Built reconciliation queries to compare source and target counts, check timestamps, and identify duplicate or missing records.
- Investigated production failures across ADF, Databricks, Snowflake, and Splunk, then handled reruns and downstream verification.
- Maintained runbooks covering recurring failure patterns, validation steps, rerun procedures, and escalation paths.
- Worked with business analysts, QA teams, and support teams to confirm data issues, validate fixes, and keep reporting cycles moving.
- **Tech:** Azure Data Factory, ADLS Gen2, Azure Databricks, PySpark, Spark SQL, Snowflake, SQL Server, Oracle, Splunk, Azure DevOps, REST APIs, Python, SQL, Git

PROJECTS

Chicago Crime Analytics Data Lake | *AWS Lambda, EventBridge, S3, Glue, Athena, QuickSight, Python, SQL* | [GitHub](#)

- Built a scheduled pipeline that pulls Chicago crime records from the Socrata API with Lambda and EventBridge, stores raw data in S3, and writes cleaned, date-partitioned Parquet for Athena analysis.
- Registered curated tables in the Glue Data Catalog and built QuickSight views for category, location, and time-of-day trends, keeping the raw layer available for replay when transformation logic changes.

Azure Databricks Lakehouse Pipeline | *Databricks, Delta Lake, DLT, Unity Catalog, PySpark*

- Built a Bronze/Silver/Gold lakehouse with ADF for ingestion, PySpark for validation and transformation, and Delta Lake for versioned, replayable storage.
- Added Delta Live Tables expectations, Kafka with Structured Streaming, and Unity Catalog access controls to cover data quality, near-real-time ingestion, and governance in the same pipeline.

Big Data Processing with Spark | *PySpark, Spark SQL, HDFS, Hive, GCP Dataproc, MapReduce* | [GitHub](#)

- Processed CSV, JSON, and XML datasets on a Dataproc Hadoop/HDFS cluster, using PySpark and Spark SQL for schema checks, cleansing, joins, and aggregations against a Hive metastore.
- Used the Spark UI to trace a shuffle-heavy join, switched the smaller side to a broadcast join, and retuned shuffle partitions to reduce unnecessary data movement and shorten the job runtime.

EDUCATION

Master of Science in Computer Science

Illinois Institute of Technology

Aug. 2023 – May 2025

Chicago, IL

Bachelor of Engineering in Computer Science

Osmania University

Aug. 2018 – April 2022

Hyderabad, India